

Work sheet 2

Due Thursday, 09. November 2017, 12:00 pm on Uniworx.

Please submit your answers **in a single PDF document** if not stated otherwise.

1. Pipelining I

Assuming we designed a MIPS processor where the pipeline stages have the following durations:

IF	= 20ns
ID	= 10ns
EX	= 20ns
MEM	= 35ns
WB	= 10ns

- What is the processor's maximum clock frequency?
- How could we achieve a performance gain in the next processor generation?
- A competitor developed a processor with 50 pipeline stages and advertises it to yield a 10x performance increase compared to our 5-stage MIPS processor.
How can this statement be argued?
How can we counter this argument?

2. Pipelining II

Write a benchmark application that allows to measure and compare throughput of scalar operations (addition, division, multiplication of integers (int, 4 bytes)) and floating point operations (double, 8 bytes) in the following two scenarios:

Scenario I: No (or infrequent) dependencies between operations so that CPU can pipeline instructions effectively.

Scenario II: Every operation depends on the result of a previous operation so that no pipelining is possible.

Document and discuss your results:

How does throughput change depending on the chosen operation (add/mul/div), data type (int/double), and scenario?
($3 \times 2 \times 2 = 12$ results)

Submit your measurement results, your source code, and the Makefile used to build the benchmark application.

Notes:

- Read the README file attached to this assignment.
- A Makefile and a C source template is provided to build your application.
You may modify the Makefile or write your own, of course.

- Avoid caching of operation results as the benchmark would otherwise measure memory access instead of ALU performance.
- Implement the benchmark application in C, assembly and/or intrinsics are not required (but viable).
- Use compiler options like `-mno-sse` for gcc to avoid generating MMX/SSE/AVX instructions and to restrict measurements to conventional scalar operations (this is already configured in the attached Makefile).

3. Data- and Instruction Processing

We discussed several concepts of parallel processing in the lecture:

- What are the conceptual differences between pipelining, vector processing and multi-threading?