

Prof. Dr. D. Kranzlmüller, Dr. K. Furlinger

Parallel Computing

WS 2017/18

Session 3: Optimizing DGEMM

Tobias Fuchs, M.Sc.

tobias.fuchs@nm.ifi.lmu.de



Today's session:

Preparation for the next assignment (online later today):

*„Optimize a given naïve implementation of
Matrix-Matrix multiplication”*

- This is a “project assignment” due in **3 weeks**.

Today's session:

Preparation for the next assignment (online later today):

*„Optimize a given naïve implementation of
Matrix-Matrix multiplication”*

- This is a “project assignment” due in **3 weeks**.
- ***It is unfathomably relevant to the final exam.***
- You may work in groups but we really don't recommend freeloading. Really learn these concepts. ***Breathe*** them.
- There will still be a new 1-week assignment next week.

Schedule:

Session 3	today	Preparation for the 3-week assignment (DGEMM)
Session 4		Discussing your progress in the DGEMM-assignment Preparation for new assignment
Session 5		Solutions to session 4 assignment

Today's session:

Preparation for the next assignment (online later today):


- Fundamentals of DGEMM
(Dense General Matrix-Matrix Multiplication)
- Cache and locality
- Other optimization techniques like loop unrolling
- Restricting to sequential variant for now

Recap: Matrices



A **matrix** is a rectangular array of values arranged in rows and columns (d'uh).


An $n \times m$ matrix has
... ?



A **matrix** is a rectangular array of values arranged in rows and columns (d'uh).

An $n \times m$ matrix has
 n rows
 m columns

Here:
 3×4 matrix



Indices in rectangular arrays

Matrix notation

$A_{ij} : (i,j)$

1,1	1,2	1,3	1,4
2,1	2,2	2,3	2,4
3,1	3,2	3,3	3,4

Cartesian notation

$A_{xy} : (x,y)$

1,1	2,1	3,1	4,1
1,2	2,2	3,2	4,2
1,3	2,3	3,3	4,3

Memory Storage Order

not to be confused with *memory ordering*,
i.e. the CPUs ability to reorder memory operations



A matrix has **two logical dimensions**.

Memory, however is linear (**one physical dimension**).

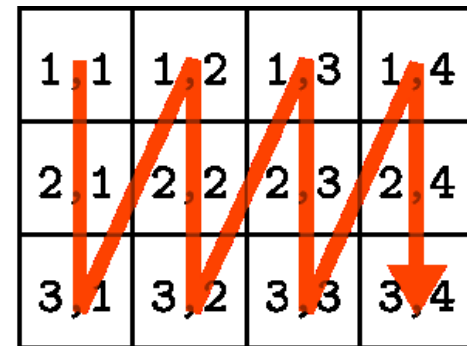
Memory order describes how multi-dimensional values are stored in linear memory.

Column-major

Index in **left-most**
dimension moves faster.

This is the Fortran style.

1,1	1,2	1,3	1,4
2,1	2,2	2,3	2,4
3,1	3,2	3,3	3,4



1,1	2,1	3,1	1,2	2,2	3,2	...
-----	-----	-----	-----	-----	-----	-----

Row-major

Index in **right-most**
dimension moves faster.

This is the C style.

1,1	1,2	1,3	1,4
2,1	2,2	2,3	2,4
3,1	3,2	3,3	3,4

1,1	1,2	1,3	1,4	2,1	2,2	...
-----	-----	-----	-----	-----	-----	-----

```
int M[N][M]  
typeof(M[0]) -> int[M]
```

i.e. $M[i]$ points to an array of row values (see C intro)

Matrix Product

Basics and Optimization



Matrix Product Basics

- Matrices are *arrays of numbers*.
- Different from elemental numbers (integer, complex, ...), there is no unique way to define “*the*” multiplication of matrices.
- “Matrix multiplication” may refer to:

Hadamard product

Kronecker product

Matrix product

entry-wise, like addition

outer product, block-matrix

the one you know from school

Matrix Product Basics

$$C = AB$$

A: $n \times m$ matrix

B: $m \times p$ matrix

C: ?

Matrix Product Basics

$$C = AB$$

A: $n \times m$ matrix

B: $m \times p$ matrix


C: $n \times p$ matrix

Matrix Product Basics

$$(AB)_{ij} = \sum_{k=1}^m A_{ik} B_{kj}$$

Notation in summation convention:

$$(AB)_{ij} = A_{ik} B_{kj}$$



Matrix Product – Naïve Implementation

```
for (index i = 0; i < n; i++) {  
    for (index j = 0; j < p; j++) {  
        double sum = 0.0;  
        for (index k = 0; k < m; k++) {  
            sum += A[i][k] * B[k][j];  
        }  
        C[i][j] = sum;  
    }  
}
```

Discuss with respect to memory layout of the matrices.

Matrix Product – Naïve Implementation

```
for (index i = 0; i < n; i++) {  
    for (index j = 0; j < p; j++) {  
        double sum = 0.0;  
        for (index k = 0; k < m; k++) {  
            sum += A[i][k] * B[k][j];  
        }  
        C[i][j] = sum;  
    }  
}
```

How about loop unrolling?

Considering Cache Locality

Arrays are contiguous memory blocks, so large chunks of them will be loaded into the cache upon first access.

$A[3][k]$	$k =$	0	1	2	3	4	5	6	7	8	first cache line
$A[4][k]$	$k =$	0	1	2	3	4	5	6	7	8	
...											
$B[6][j]$	$j =$	0	1	2	3	4	5	6	7	8	last cache line

Considering Cache Locality

- Cache locality matters a **lot**.
- Loading data from main memory into cache takes **hundreds** of CPU cycles
- Cache misses dominate running time more than the actual calculations
- In an upcoming session, I will present tools to actually measure performance, e.g. using cache miss counters.

Lessons Learned for Optimization

- Make yourself familiar with blocking / tiling optimization techniques (you will find lots of references in the web).
- Worry about cache first, then go for `-O` flags, loop unrolling etc.

Last words

- Have fun with performance tweaking!
Performance optimization is all about “cheating”.
- We will discuss your questions and progress next week.
- **Do not hesitate** to contact me when you’re stuck.
The best coders became champions because they dared to ask lots of “stupid” questions in their lives.

Tobias Fuchs

tobias.fuchs@nm.ifi.lmu.de

www.mnm-team.org/~fuchst

